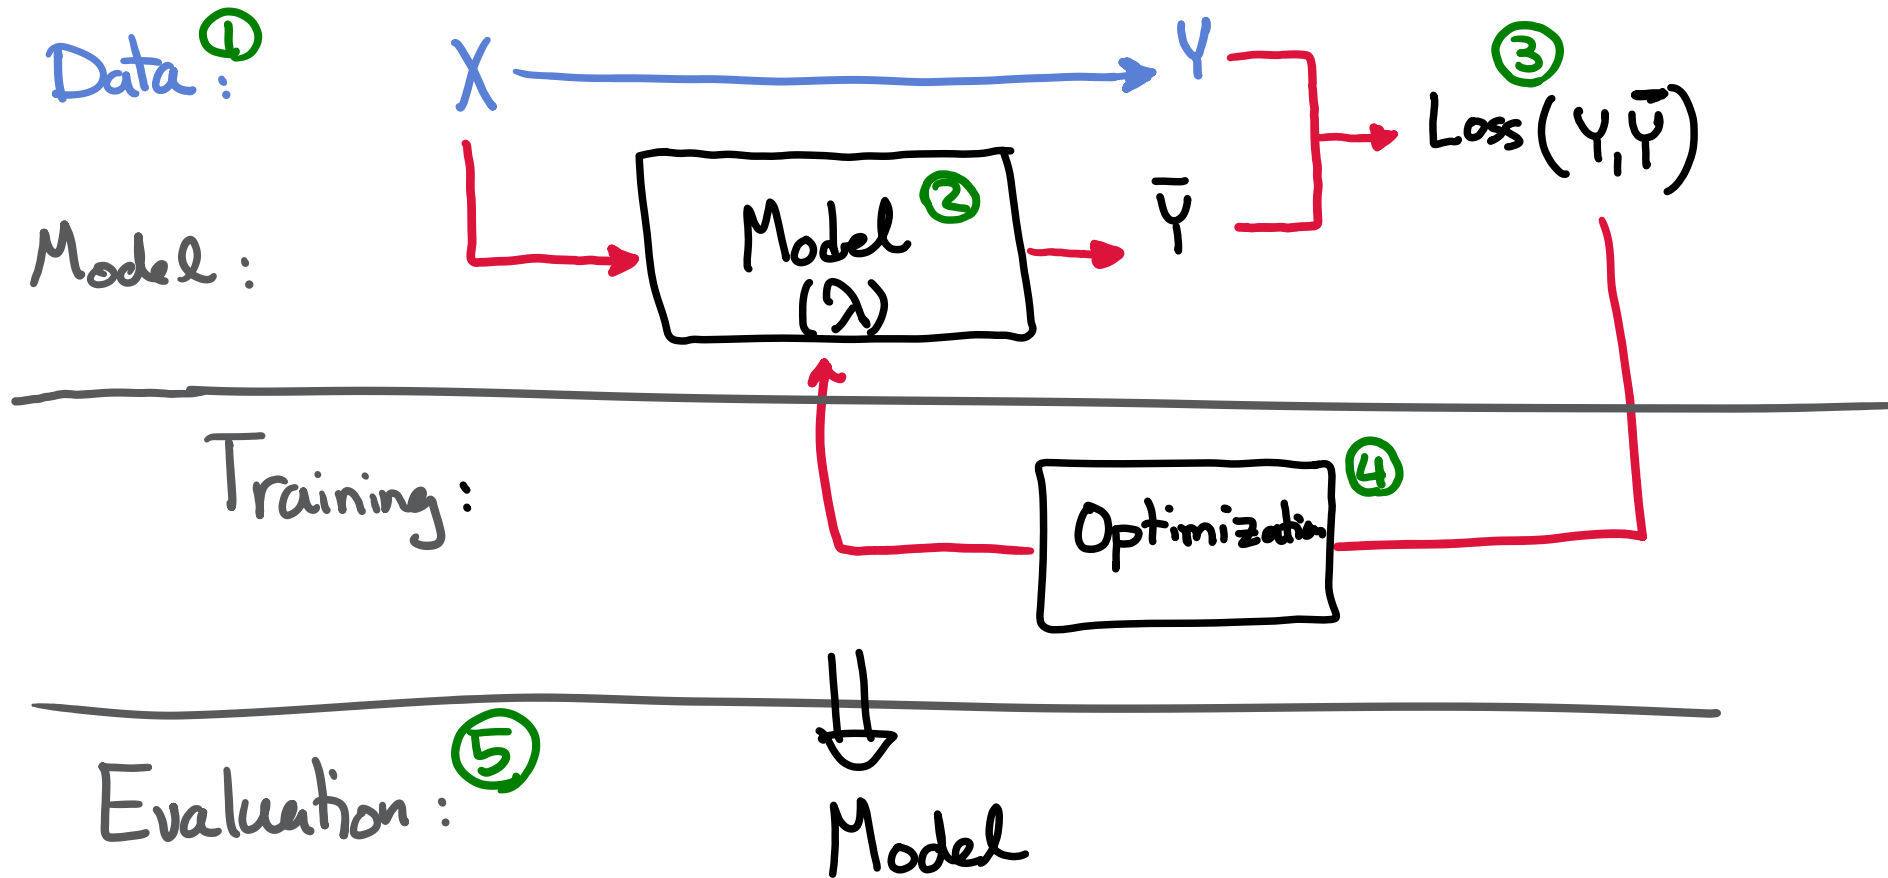
The background of the slide features a complex, abstract pattern of blue and white wavy lines, resembling a topographical map or a fluid simulation, set against a solid black background. The lines are dense and create a sense of depth and movement.

# Machine Learning in Physics: **Model Evaluation**

Sadegh Raeisi

# Supervised: Ingredients



# Outline

What's a good model?


Bias and Variance

Metrics


Model Tuning

# A good model

# A good fit vs a good model




# A good fit vs a good model




# A good fit vs a good model


Linear



Quadratic




Higher order




What's the problem?

How can we solve it?






# Good fit vs Good Prediction



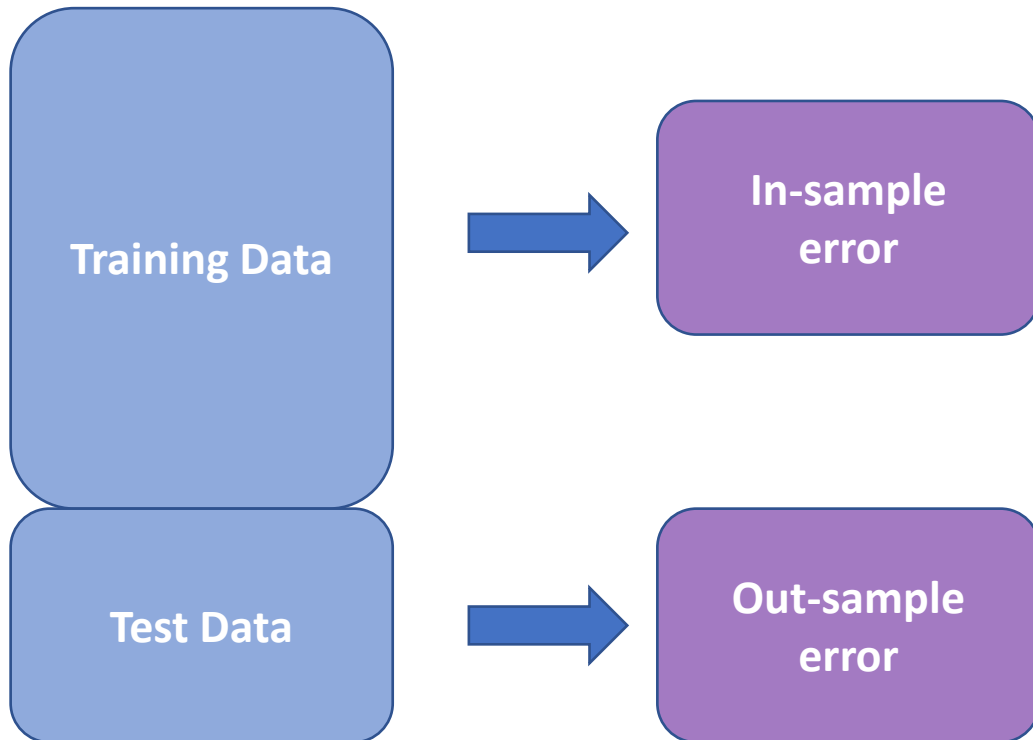
# How can check the prediction power of a model?

- In-sample vs out sample error



# Bias vs Variance

# Paper




What do they depend on?

How can we reduce them?

Which one do we care more about?

# Bias and Variance



What do they depend on?

How can we reduce them?

Which one do we care more about?


# Bias and Variance: more detail

$$\mathcal{L}(Y, \bar{Y}) = \sum_i \left( Y^i - f_w(X^i) \right)^2$$

This depends  
on the Data.


$$\mathcal{L}_{D_i}(Y, \bar{Y})$$

All the possible data



$$\mathbb{E}_D[\mathcal{L}(Y, \bar{Y})]$$

# Bias and Variance: more detail



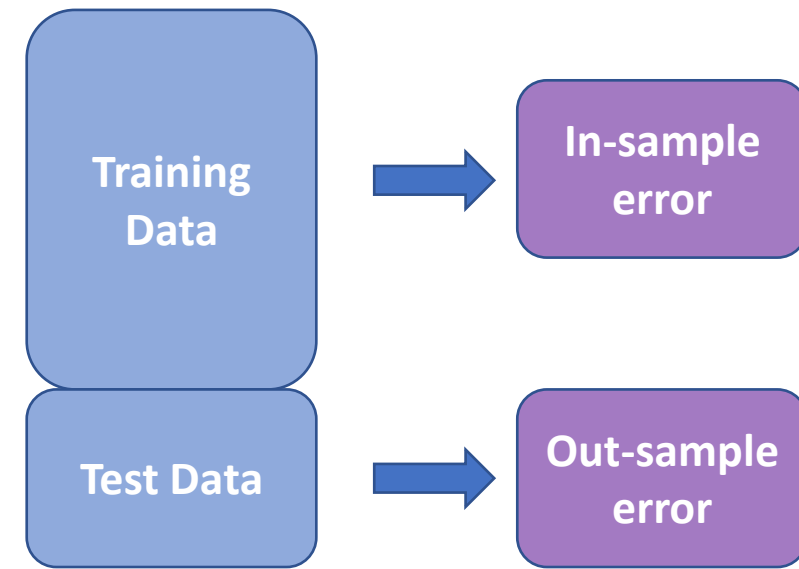
$$\mathbb{E}_D[\mathcal{L}(Y, \bar{Y})]$$

$$= \sum_i \mathbb{E}_D \left( Y^i - \mathbb{E}_D \left( f_w(X^i) \right) + \mathbb{E}_D \left( f_w(X^i) \right) - f_w(X^i) \right)^2$$

$$= \underbrace{\sum_i \left( Y^i - \mathbb{E}_D \left( f_w(X^i) \right) \right)^2}_{\text{Bias}^2} + \underbrace{\sum_i \mathbb{E}_D \left( \mathbb{E}_D \left( f_w(X^i) \right) - f_w(X^i) \right)^2}_{\text{Variance}}$$



# For a good model

$$\mathbb{E}_D[\mathcal{L}(Y, \bar{Y})]$$



$$= \sum_i \left( Y^i - \mathbb{E}_D \left( f_w(X^i) \right) \right)^2 + \sum_i \mathbb{E}_D \left( \mathbb{E}_D \left( f_w(X^i) \right) - f_w(X^i) \right)^2$$

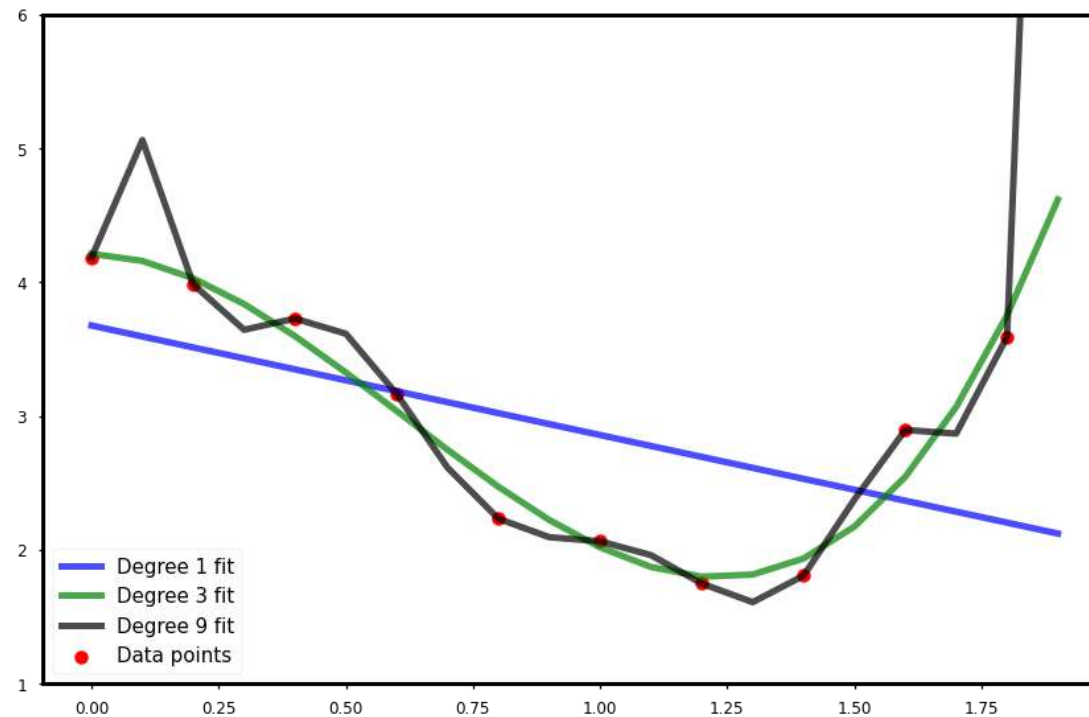

**Bias<sup>2</sup>**                      **Variance**






How can we reduce bias?

How can we reduce variance?




Check the code!

# Validation curve




# How much complexity?




# How much data?

Do we have enough data?



# Learning curve



# Regularization

$$\mathcal{L}(Y, \bar{Y}) = \sum_i \left( Y^i - f_w(X^i) \right)^2 + \alpha \|w\|_l$$

$$\mathbf{L2}: \|w\|_2 = \sum_i w_i^2$$

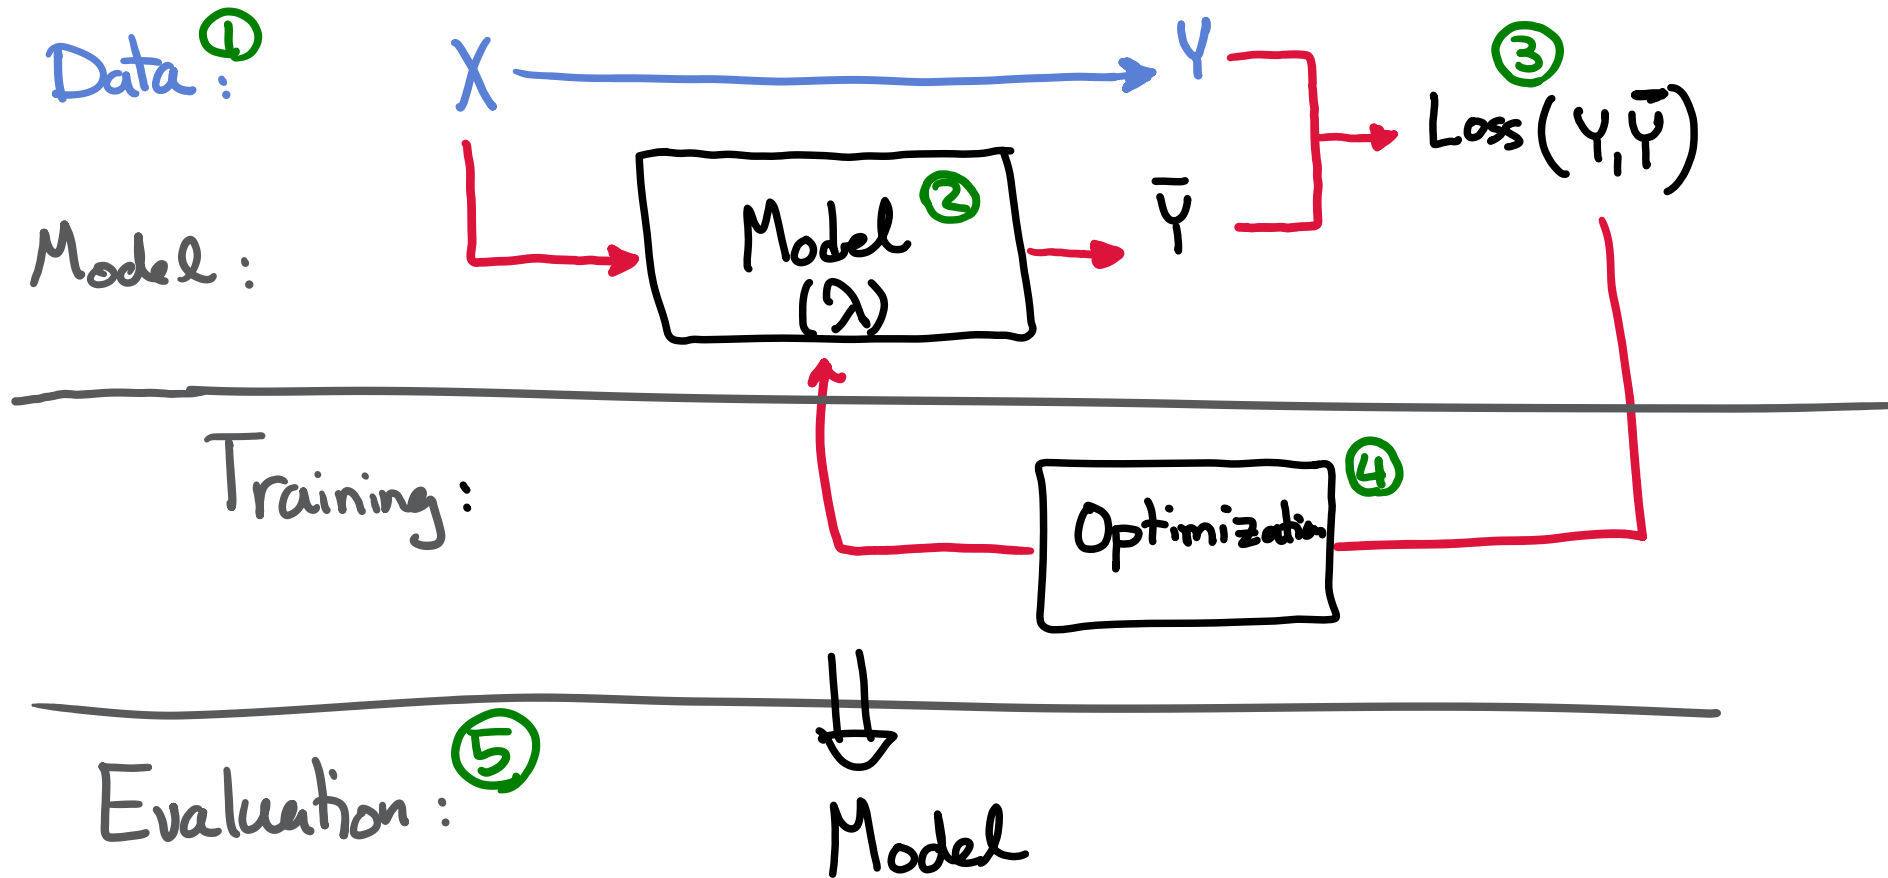
$$\mathbf{L1}: \|w\|_1 = \sum_i |w_i|$$

# Metrics




# Recap

# Supervised: Ingredients

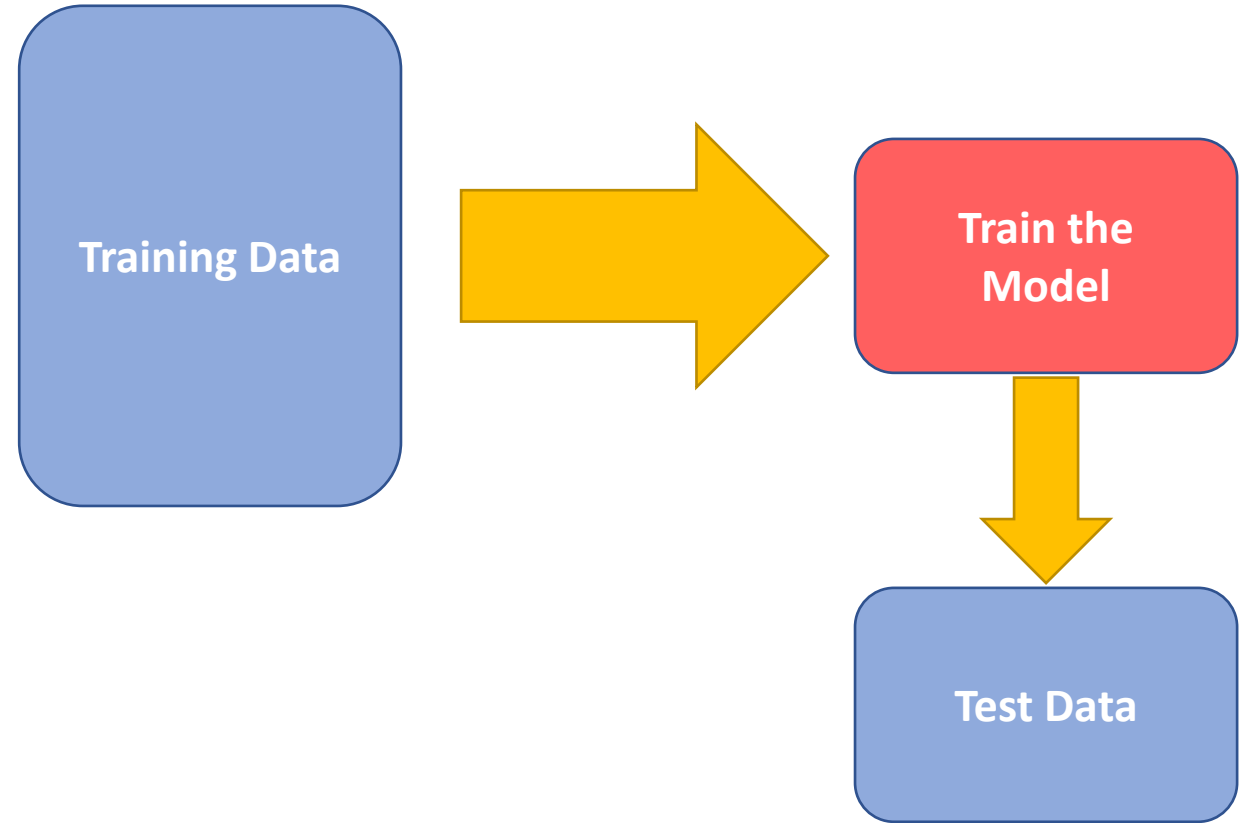


# Code



```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X , Y, random_state=0)
```

# Code



```
from sklearn.linear_model import SGDClassifier
```

```
clf = SGDClassifier()  
clf.fit(X_train, Y_train)
```

```
y_predict = clf.predict(X_test)  
error = np.abs(Y_test - y_predict).sum() / len(Y_test)
```

# Code: full pipeline

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X , Y, random_state=0)
```

```
from sklearn.linear_model import SGDClassifier
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures

### Training the model
clf_pipeline= Pipeline([('scaler', StandardScaler() ),
                        ('p_transformer', PolynomialFeatures(degree = 3)),
                        ('clf', SGDClassifier())])
clf_pipeline.fit(X_train,Y_train)
```

```
### Testing the model
y_predict = clf_pipeline.predict(X_test)
out_error = np.abs(Y_test - y_predict).sum() / len(Y_test)
in_error = np.abs(Y_train - clf_pipeline.predict(X_train) ).sum() / len(Y_train)
```